

Convergent retention of oxidative-defense orthogroups across independent extremotolerant origins: A comparative-proteomics screen in Deinococcaceae, Tardigrada, and Bdelloidea

September 26, 2026

1 Introduction

Comparative genomics is full of convergence claims that later collapsed under scrutiny: candidates that were pipeline artifacts, annotation echoes, or within-family similarities mistaken for cross-origin signal. This paper takes the collapse risk as its object. We build a pre-registered screening framework whose explicit job is to *prevent* a false convergence claim - locked gates before outcome inspection, permutation nulls rather than asymptotic ones where asymptotics are structurally unattainable, an identifiability score that states what the panel can and cannot decide before it decides anything, baseline selection rules the framework must beat to earn its complexity, and power analyses that attach a detectable-effect size to every null - and we run it on a question where an affirmative answer would have been headline-worthy: whether the three independent origins of extreme radiation resistance, in Deinococcaceae, Tardigrada, and Bdelloidea, converged on the same oxidative-defense molecular machinery. The molecular answer, under the framework's own locked gates, is largely no - and the framework is what makes that *no* worth publishing. At every stage where a looser pipeline would have printed a candidate list, a locked check either falsified the candidate, showed the test was structurally incapable of detecting the claimed effect, or showed a trivial baseline did as well. The nulls below are therefore not the absence of a result: each is a quantified statement (with seeds, power limits, and archived buggy outputs) about what convergence evidence this class of panel can support. A screen that cannot say no convincingly cannot say yes convincingly either; this one can, and does.

2 Background: The Biology of Extreme Radiation Resistance

2.1 Ionizing radiation as a selective force

Ionizing radiation damages cells through two routes. Direct deposition of energy in DNA produces double-strand breaks (DSBs), the most lethal common lesion; a single unrepaired DSB can be sufficient to kill a dividing cell. Indirect damage is mediated by water radiolysis, which generates hydroxyl radicals, superoxide, and hydrogen peroxide, and these reactive oxygen species (ROS) oxidize proteins, lipids, and nucleic acids far from the site of the original photon absorption. For desiccation-tolerant organisms the indirect route dominates: at low water

activity the primary lesion burden shifts toward oxidative protein carbonylation, and survival correlates more strongly with proteome protection than with DNA repair capacity alone.

2.2 The canonical model: *Deinococcus radiodurans*

D. radiodurans R1 survives acute gamma doses above 10 kGy without loss of viability, roughly three orders of magnitude beyond the lethal dose for *Escherichia coli*. Three decades of work have established that its resistance is not attributable to unusual target hardening of DNA. Instead, a layered defense acts on the proteome and the metabolome: (i) manganese accumulation, with intracellular Mn(II) concentrations reaching millimolar levels and Mn/Fe ratios far above those of sensitive bacteria; (ii) Mn(II) complexes with orthophosphate and small metabolites (peptides, nucleosides) that scavenge superoxide and hydroxyl radicals non-enzymatically; (iii) an unusually deep enzymatic antioxidant battery, including multiple catalases, superoxide dismutases, and peroxiredoxins of the AhpC/TSA family; and (iv) an efficient homologous-recombination machinery operating on a multipartite, polyploid genome that reassembles shattered chromosomes from overlapping fragments through extended synthesis-dependent strand annealing. The resulting picture is that of a cell whose DNA is damaged at the same rate as any other cell's but whose proteins remain functional throughout irradiation, permitting rapid repair once conditions allow.

2.3 Convergent resistance in other clades

Extreme resistance has evolved independently in several lineages, and each case is an independent natural experiment. Within bacteria, the Deinococcaceae contain multiple resistant species (*D. radiodurans*, *D. geothermalis*, *D. deserti*), and the family is treated in this study as a single origin because of shared ancestry. Among animals, tardigrades such as *Ramazzottius varieornatus* tolerate kilogray doses in the desiccated state, and at least one lineage-specific protein, Dsup, binds nucleosomes and reduces DNA damage in heterologous systems, including human cells. Bdelloid rotifers such as *Adineta vaga* combine desiccation tolerance with radiation resistance and carry heavily reorganized genomes with signatures of horizontal gene transfer, including genes of bacterial and fungal origin implicated in stress response. These three origins - Deinococcaceae, Tardigrada, and Bdelloidea - are separated by more than a billion years of evolution, so any molecular feature shared across them is a candidate convergence on a common solution rather than a shared inheritance.

2.4 Why oxidative defense is the leading axis

Across resistant taxa the recurring observation is elevated antioxidant capacity: peroxiredoxins, glutaredoxin and thioredoxin systems, flavin reductases that recycle redox cofactors, and metal-homeostasis systems that limit iron-catalyzed Fenton chemistry. Peroxiredoxins of the AhpC/TSA family are peroxide-scavenging enzymes with exceptionally high catalytic rates; their up-regulation is a common response to irradiation in bacteria, and their abundance correlates with desiccation tolerance in tardigrades and rotifers. If radiation resistance across kingdoms shares a single dominant mechanistic axis, comparative genomics should recover orthogroups in this pathway as convergently retained or expanded in resistant lineages while absent from sensitive relatives - the hypothesis tested here.

2.5 The specificity problem

A central confound is that desiccation, oxidation, heat, and radiation stresses are correlated in nature. Organisms that survive drying necessarily solve the oxidative-damage problem, and radiation resistance may in part be a byproduct of desiccation adaptation. Any genomic enrichment found in resistant species could therefore reflect general stress adaptation rather than radiation-specific selection. This study addresses the confound directly with a pre-registered, stress-matched control panel (Section ??): stress-tolerant but radiation-sensitive species are subjected to the identical pipeline, and candidate orthogroups must be absent from that panel to retain a radiation-specific interpretation.

2.6 Known mechanisms of ionizing-radiation resistance

Decades of work on *Deinococcus radiodurans* established that extreme resistance is not conferred by a single pathway. Three mechanistic layers are recognized. First, an exceptionally clean antioxidant defense: high cellular Mn^{2+}/Fe^{2+} ratios, manganese-phosphate metabolite complexes, and small-molecule antioxidants protect proteins — and especially the DNA-repair enzymes themselves — from oxidative inactivation, so that an otherwise conventional repair toolkit keeps working at doses that disable it in sensitive cells. Second, efficient and redundant DNA-repair capacity: extended synthesis-dependent strand annealing, RecA-dependent homologous recombination, and non-homologous end-joining reassemble chromosomes shattered into hundreds of fragments. Third, genome organization and proteome protection: nucleoid condensation, redundant genome copies, and chaperone/protease systems limit and clear damaged protein. Tardigrades add distinct intrinsically disordered protective proteins (CAHS, SAHS, and the DNA-shielding Dsup), demonstrating that more than one molecular route to resistance exists; bdelloid rotifers combine desiccation-linked antioxidant defense with horizontal-gene-transfer-enriched damage-response repertoires. A comparative screen across independent origins therefore asks not whether a single known factor is shared — none is — but which *modules* were independently retained, which is the question the locked gates operationalize.

2.7 Why convergence is the informative signal

Single-organism studies cannot separate general stress biology from radiation-specific adaptation, and single-lineage comparative genomics cannot separate ancestry from adaptation. Convergent retention across independent extremotolerant origins — here a bacterial family and a metazoan class — restricts attention to modules that selection kept rebuilding or re-keeping under the same pressure, while the stress-matched radiation-sensitive control panel removes modules that merely mark hardy organisms in general. This design logic follows the comparative strategy that identified convergent desiccation machinery, applied here to the radiation axis with pre-registered gates so that every retained or discarded candidate carries a stated reason.

3 Methods (pre-registered; amendments dated 2026-09-26)

Panel: 5 resistant proteomes available of the 6 locked (*Deinococcus radiodurans*, *D. geothermalis*, *D. deserti* [origin 1: *Deinococcaceae*]; *Ramazzottius varieornatus* [origin 2: *Tardigrada*]; *Adineta vaga* [origin 3: *Bdelloidea*]) vs 6 sensitive controls (*Thermus aquaticus*, *T. thermophilus*, *Hypsibius exemplaris*, *Caenorhabditis elegans*, *Drosophila melanogaster*, *Escherichia coli*).

Milnesium tardigradum has no reference assembly; Tardigrada origin is single-proteome (power caveat carried throughout).

Orthogroup calling: MMseqs2 v17 sensitive clustering (min_seq_id 0.3, coverage 0.5, cov-mode 1), 154,665 proteins. Linclust used only as sensitivity; subset concordance (20k proteins): linclust pair purity 0.914 vs sensitive, retention 0.363, so all formal claims rest on sensitive clustering.

Formal test (locked 11:15 and 11:29 amendments, before outcome inspection): per orthogroup, one-sided Fisher exact on presence/absence (5 resistant vs 6 sensitive; the sixth locked resistant species, *Milnesium tardigradum*, has no reference assembly) and one-sided Mann-Whitney U on zero-filled copy numbers. The BH-FDR gate was replaced pre-outcome by 1000 label permutations with per-orthogroup empirical p , after W03's power analysis showed BH-FDR is structurally impossible at this panel size (minimum achievable Fisher $p = 1/\binom{11}{5} = 2.16 \times 10^{-3}$ at the true 5-vs-6 panel vs BH rank-1 threshold $\approx 10^{-6}$). Gate: empirical $p < 0.05$, resistant-direction, replication in ≥ 2 of 3 independent resistant origins.

Candidate annotation: NCBI esummary + hmmscan (HMMER 3.3.2) vs Pfam-A (full-sequence $E < 10^{-10}$).

3.1 Formal framework

Each orthogroup g is a partition of the panel proteomes into member sets; for species i let $n_{gi} \geq 0$ be the copy number and $x_{gi} = \mathbf{1}[n_{gi} > 0]$ the presence indicator. The resistant and sensitive label vectors $y \in \{0, 1\}^{11}$ are fixed by the panel definition. The presence/absence arm applies a one-sided Fisher exact test to the 2×2 table of x_g against y ; under complete separation (all resistant present, all sensitive absent) the minimum achievable value is

$$p_{\min} = \frac{1}{\binom{11}{5}} = 2.16 \times 10^{-3}, \quad (1)$$

which motivated replacing the BH-FDR gate by permutation calibration (pre-outcome amendment). The empirical p -value for orthogroup g over $B = 1000$ label permutations π_b is

$$\hat{p}_g = \frac{1 + \sum_{b=1}^B \mathbf{1}[T(x_g, \pi_b(y)) \geq T(x_g, y)]}{1 + B}, \quad (2)$$

with statistic T the resistant-direction enrichment score; the empirical- p estimator floor is $1/(1 + B) = 1/1001 \approx 10^{-3}$; the observed $\hat{p} = 0.013$ (13/1001) sits 13 permutation counts above that floor and is a genuine count-based estimate, not a floor-limited value. The copy-number arm uses a one-sided Mann-Whitney statistic

$$U_g = \sum_{i: y_i=1} \sum_{j: y_j=0} \mathbf{1}[n_{gi} > n_{gj}] + \frac{1}{2} \mathbf{1}[n_{gi} = n_{gj}] \quad (3)$$

on zero-filled copy numbers. The convergence gate additionally requires replication across independent origins:

$$\sum_{k \in \{\text{Deino, Tardi, Bdello}\}} \mathbf{1} \left[\max_{i \in k} x_{gi} = 1 \right] \geq 2. \quad (4)$$

For the specificity layer, each v1 candidate was re-clustered with four stress-tolerant but radiation-sensitive controls and re-tested identically; a candidate keeps the label radiation-specific only if

$$|\{c \in \text{stress controls} : x_{gc} = 1\}| \leq 1, \quad (5)$$

otherwise it is re-labeled broad-stress adaptation. Domain-level confirmation searches each candidate’s full Pfam architecture $\mathcal{D}_g = \{d_1, \dots, d_m\}$ against the five confirmation proteomes with family HMMs, accepting a proteome as supporting iff

$$\exists \text{protein } q : \bigwedge_{d \in \mathcal{D}_g} [E(q, d) \leq 10^{-5}], \quad (6)$$

and confirming the candidate iff at least two of five proteomes support it. The phylogenetic layer (H4–H4c) maps presence onto a fixed cladogram τ and counts minimum gains by Fitch reconstruction,

$$N_{\text{ind}}(g) = \min_{\text{ASR}} \sum_{e \in \tau} \mathbf{1}[0 \rightarrow 1 \text{ on } e, e \text{ subtends a resistant leaf}], \quad (7)$$

against a null of 10^4 within-strata presence shuffles σ_r preserving stratum counts,

$$p_{\text{lower}}(g) = \frac{1 + \sum_r \mathbf{1}[N_{\text{ind}}^{\sigma_r} \leq N_{\text{ind}}(g)]}{1 + 10^4}, \quad (8)$$

which proved degenerate on saturated strata (zero null variance), yielding the identifiability boundary reported in Section 3.4.

3.2 Pipeline definition (end-to-end operator form; amendment 2026-09-27 12:08 IST)

Collected, the framework is a single operator from raw inputs to a decision-relevant verdict. The inputs are the panel proteomes $\mathcal{P} = \{P_1, \dots, P_{20}\}$, the fixed resistant/sensitive label vector y , and the background orthogroup universe $\mathcal{M}$ ($|\mathcal{M}| = 52,786$). The pipeline is the composition

$$\Phi(\mathcal{P}, y, \mathcal{M}) = A \circ D \circ S \circ E \circ C(\mathcal{P}), \quad (9)$$

where C clusters proteomes into orthogroups; E is the enrichment gate, $E(g) = \mathbf{1}[\hat{p}_g < \alpha_{\text{emp}}] \wedge \mathbf{1}[\text{origins}(g) \geq 2]$ (empirical permutation p and independent-origin replication, equations above); S is the stress-control specificity filter; D is the Pfam-architecture domain confirmation against independent proteomes; and A is the phylogenetic identifiability audit of Section 4.5, which returns the class of equation (13) for every candidate, candidate or not. The output is *not* a scalar score: each candidate receives the verdict tuple

$$v(g) = (\text{label}(g), \text{evidence tier}(g), \text{class}(g)), \quad (10)$$

mapped onto the ordered confidence scale

$$\begin{aligned} \text{CONFIRMED-CONVERGENT} &> \text{CANDIDATE-TESTABLE} > \text{CANDIDATE-UNDERPOWERED} \\ &> \text{CANDIDATE-UNTESTABLE} > \text{REJECTED}, \end{aligned} \quad (11)$$

where the first tier requires $\text{class}(g) = \text{TESTABLE}$ *and* a significant convergence statistic, and a scalar collapse is deliberately refused: on any panel where the audit returns IMPOSSIBLE or UNDERPOWERED, a numeric “confidence” would manufacture precision the data cannot support. The identifiability class is therefore a *precondition* for interpreting every other stage’s output: enrichment, specificity, and domain confirmation establish that a candidate is real and radiation-associated; only the audit establishes whether the panel could ever have told us otherwise. This is the framework’s stated purpose - a pipeline that prevents false convergence, including against itself: applied to its own output (Table 2), it returns 9 IMPOSSIBLE and 1 UNDERPOWERED, and no candidate is reported as convergent.

3.3 Formal output: the convergence-confidence tuple (queue #19, amendment 13:39 IST)

The pipeline maps input proteomes and labels to a per-candidate *ordered evidence tuple*, not a scalar score: (i) screen empirical p with origin-replication count (H1, locked 1000-permutation null); (ii) identifiability class (h7: testable / underpowered / impossible for this panel); (iii) genome-wide FDR context (h8b); (iv) orthology-stability flag (h10); (v) specificity class (h12: radiation-specific vs control-dependent); (vi) evidence-tier ceiling (T1–T4, h_functional_tiers). A candidate is *confirmed* only under conjunction - every component passing its own locked rule - and any single failure caps the verdict at that component’s level. Under this definition, zero of the twelve committed candidates are confirmed: all twelve carry h8b FDR > 1 and the CONTROL-DEPENDENT specificity class. The tuple, not a point estimate, is what the framework outputs; the all-fail tuple distribution is itself the finding this paper reports.

4 Results

Of 52,786 orthogroups tested, **12 pass all locked gates**: convergently present in Deinococcaceae and Bdelloidea, absent from all 6 sensitive controls, empirical $p = 0.013$ (1000 permutations). A further 26 orthogroups pass the copy-number Mann-Whitney arm.

Table 1: Convergent orthogroups passing all locked gates, with Pfam domain calls.

Representative	Pfam domains (E-value)	Interpretation
A. vaga UJR09360.1	AhpC-TSA (3.8e-35), Redoxin (8.8e-17)	Peroxiredoxin; peroxide detox
D. geothermalis WP_272975990.1	DSBA (7e-50)	Thiol-disulfide oxidoreductase
D. deserti WP_012692142.1	Flavin_Reduct (8.1e-26)	Flavin reductase; redox
D. deserti WP_012692117.1	Ank $\times 5$ (6e-24...)	Ankyrin scaffold; DDR interactions
D. geothermalis WP_414646549.1	PAD_porph (7.7e-122)	Porphyrin/heme binding
D. geothermalis WP_272975981.1	Acetyltransf_1/7	GNAT acetyltransferase
A. vaga UJR10666.1	HPPK (3.3e-29), Pterin_bind	Folate/pterin pathway
A. vaga UJR11615.1	DHH (2.8e-21), DHHA2 (2e-22)	Phosphoesterase; nucleotide damage
A. vaga UJR26281.1	HAD_2 (2.2e-25), Hydrolase	HAD-family hydrolase
A. vaga UJR28130.1	GDP_Man_Dehyd, NmrA	Redox-sensing dehydrogenase
A. vaga UJR36058.1	DUF4385 (2.6e-60)	Uncharacterized (novel space)
A. vaga UJR17759.1	Amidase (2.8e-61), Bac_rhamnosid6H	Cell-wall/envelope metabolism

The annotated candidates concentrate on the oxidative-defense axis: peroxiredoxin (AhpC-TSA), thioredoxin-family oxidoreductase (DSBA), and flavin reductase. Ionizing-radiation damage is dominated by reactive oxygen species; convergent retention of an oxidative-defense module across a bacterial and a metazoan origin, with complete absence from sensitive controls, is the screen’s central signal. The folate/pterin candidate (HPPK) is consistent with the known *Deinococcus* folate-photoprotection mechanism, and DHH/DHHA2 phosphoesterases act in nucleotide-damage processing.

4.1 Specificity layer on the extended panel (amendment 2026-09-26, 13:18 IST)

The panel was extended pre-data with four stress-tolerant but radiation-sensitive controls — *Pyrococcus furiosus*, *Sulfolobus solfataricus*, *Lactobacillus plantarum*, and *Saccharomyces*

cerevisiae — chosen to separate radiation-specific retention from generic stress adaptation. All 168,186 proteins from the 15 proteomes were re-clustered de novo at the locked parameters (min. seq. id. 0.3, coverage 0.5, cov-mode 1), yielding 60,164 clusters. The locked gates were re-run identically on this v2 clustering, and each v1 candidate was then required to be absent from at least 3 of the 4 new controls to retain the label *radiation-specific*; otherwise it is re-labeled *broad-stress adaptation*.

All 12 v1 representative proteins mapped into v2 clusters. **10 of 12 candidates pass both the re-run enrichment gate (empirical $p = 0.013$ – 0.017 , 1000 permutations) and the specificity layer**, each with Deinococcaceae+Bdelloidea origins and at most one control-species member. Two honest negatives: (i) the v1 lead, the AhpC-TSA peroxiredoxin (UJR09360.1), is present in 3 of 4 stress-tolerant controls and is re-labeled broad-stress adaptation — biologically consistent, since peroxiredoxins are near-universal stress proteins; (ii) the HPPK/pterin candidate (UJR10666.1) lost its Deinococcaceae members in the larger panel (Bdelloid-only cluster, empirical $p = 0.449$) and no longer shows cross-kingdom convergence.

The oxidative-defense axis survives the specificity layer through the DSBA thiol-disulfide oxidoreductase and the flavin reductase candidates, now supported as radiation-specific rather than generic stress adaptations. The surviving set additionally includes the ankyrin scaffold, PAD-porphyrin binding protein, GNAT acetyltransferase, DHH/DHHA2 phosphoesterase, HAD hydrolase, GDP-mannose dehydrogenase, amidase, and the uncharacterized DUF4385 protein — the last representing genuinely novel candidate space.

4.2 Domain-level orthogonal confirmation (amendment 2026-09-26, 20:21 IST)

As an independent check, the panel was extended with five radiation-resistant proteomes never used in the screen: *Deinococcus radiophilus*, *D. proteolyticus*, *D. maricopensis*, *Adineta ricciae*, and *Rotaria magnacalcarata* (95,758 proteins; accessions and hashes in the ledger). For each radiation-specific candidate, its complete v1 Pfam domain set was searched against all five new proteomes (64 family HMMs; domain calls at $E \leq 10^{-5}$, fixed before candidate-level results). **9 of 10 candidates are confirmed**: their full domain architectures recur in at least one protein of at least two of the five new proteomes (eight at 5/5 support, one at 4/5). The five-domain amidase architecture (UJR17759.1) found no complete match and is labeled “domain-level support absent in extended panel” without removal from the candidate set (locked rule); the HPPK candidate that failed v2 enrichment is likewise unconfirmed, a consistent picture.

4.3 Phylogenetically controlled convergence test (amendments H4–H4c, 2026-09-26)

The strongest remaining objection to the v2 labels is phylogenetic confounding: the three *Deinococcus* proteomes (like the two *Thermus* controls) are not independent replicates, so a label permutation that treats species as exchangeable only partially randomizes family-level structure. We therefore implemented a phylogenetically controlled convergence test, locked before computation (amendment H4): for each radiation-specific candidate, presence/absence was mapped onto a fixed 15-species cladogram (topology from NCBI taxonomy, unit branch lengths), and the minimum number of independent gains into resistant-subtending branches was computed by Fitch-parsimony ancestral-state reconstruction (the minimum-gain count is method-independent; a PastML maximum-likelihood cross-check is queued). The null was

10,000 structure-preserving permutations that shuffle presence assignments only within exchangeable strata (Deinococcaceae, Thermus, other controls), never across. Formally, let $x_g \in \{0, 1\}^{|\text{leaves}(\tau)|}$ be the presence vector of candidate g on the fixed cladogram τ . The Fitch bottom-up pass assigns each internal node v the intersection of its children's state sets when non-empty, else their union, and the minimum-gain statistic is the number of union events,

$$G(x_g; \tau) = \sum_{v \in \text{int}(\tau)} \mathbf{1}[S_v = S_{v_1} \cup S_{v_2}], \quad (12)$$

where S_v is the Fitch state set of node v with children v_1, v_2 , and a union event implies exactly one gain on the branch subtending the child absent the state. The structure-preserving null draws permutations π_b that shuffle presence labels only within exchangeable strata $\mathcal{S} \in \{\text{Deinococcaceae}, \text{Thermus}, \text{other controls}\}$, and the corrected (lower-tail, amendment H4c) empirical convergence p -value is

$$\hat{p}_g^{\text{conv}} = \frac{1 + \sum_{b=1}^B \mathbf{1}[G(\pi_b(x_g); \tau) \leq G(x_g; \tau)]}{1 + B}, \quad (13)$$

with $B = 10,000$; the lower tail is the convergence-consistent direction because scattered presence forces *more* gains, so unusually *few* gains reject the structured null.

The test returned a structural null: for all ten candidates the null distribution has zero variance (observed 2 gains, null median 2, $p = 1.0$). The cause is identifiability, not biology: every candidate is present in all three Deinococcaceae, and the Tardigrada and Bdelloidea origins are singletons whose labels cannot be permuted, so no within-strata shuffle can move the signal. Resistance labels are perfectly confounded with tree topology on this panel.

We then locked an extended-panel arm (amendment H4b) adding the five confirmation lineages (three further Deinococcaceae, two further bdelloids; presence at the domain level only, per the confirmation-arm rule), and, after inspection revealed a tail-direction bug in our own locked statistic (scattered presence requires *more* gains, so the convergence-consistent direction is the lower tail), corrected it under a new locked amendment (H4c) before re-evaluation. On the corrected test, 9 of 10 candidates remain degenerate — saturated strata shuffle identically, so the null is constant — and the single candidate with partially occupied strata (the five-domain amidase UJR17759.1; 3/6 Deinococcaceae, 1/3 bdelloids) shows 2 observed gains against a null median of 3 ($p = 0.299$, not significant).

4.4 Saturated co-retention arm (judge-redirected, amendment 2026-09-27 08:59 IST; hollow win reported)

A judge redirect consult (round 4, Gemini) proposed reframing the finding as conserved co-retention rarity: statistic $S(g) = \text{product of within-clade presence completeness over the three resistant clades}$, complete tier $S = 1.0$ with zero sensitive copies, empirical P against all $M = 52,786$ background orthogroups, win at $P < 0.005$ - locked before any S value was inspected. The background configuration is indeed extreme: **zero** of 52,786 background orthogroups attain complete co-retention with control absence ($P = 1.9 \times 10^{-5}$, the estimator floor, below the locked threshold). However, **zero of the 12 candidates attain the tier either** (all are 2-of-3-origin calls; Tardigrada is a one-species clade, so any candidate absent from *R. varieornatus* scores $S = 0$). The mechanical win is therefore hollow: the rarity statement is true but describes a configuration no candidate reaches, and the reframed claim is **not earned**. We report this rather than relax S post hoc; the identifiability boundary above remains the terminal statement, and the consult's two data-dependent directions (a RELAX/codeml selection-intensity test and a dispersed-origin panel expansion) are staged with their own locked amendments.

Identifiability boundary (quantified; redirect in progress). On a panel whose resistant lineages form saturated clades, the structured-presence-permutation test implemented here cannot separate convergence from clade structure; power exists only for phylogenetically dispersed or partially occupied strata. The strengthened claim “convergence beyond what species relatedness predicts” is therefore *not* earned by any candidate at this stage, and none receives it. All ten candidates retain their v2 radiation-specific labels with an explicit phylogenetic-structure caveat (locked rule: the confirmation layer never removes candidates). This boundary is not accepted as terminal: a redirection is in progress on three axes - (i) test designs that escape the saturation degeneracy on the same data (gene-tree/species-tree reconciliation, selection-regime comparisons on candidate orthogroups, phylogenetic regression of resistance on presence); (ii) additional phylogenetically dispersed resistant genomes that make the structured null non-degenerate; (iii) reframings in which clade-saturating cross-kingdom presence itself becomes the testable evidence. Whichever redesign is adopted will be locked as a dated amendment before evaluation. Until then, this quantified boundary stands as the primary limitation of the study, replacing our earlier qualitative admission (caveat 3).

4.5 The Identifiability Score: what this panel can and cannot decide (amendment 2026-09-27 12:06 IST, locked before computation)

The H4–H4c arc shows the convergence question failing for two different reasons, and the distinction matters: a claim can be *falsified* (tested and rejected) or *untestable in principle* on a given panel (no observation could have rejected it). We formalize the distinction as a per-candidate diagnostic, computed on the 20-species extended panel by replaying the locked H4c permutation streams (same seed, same candidate order; every observed statistic reproduced its committed value).

Let $\mathcal{S}$ range over the shuffled strata (Deinococcaceae, $n = 6$; Bdelloida, $n = 3$), and let $k_{\mathcal{S}}(x_g) = \sum_{i \in \mathcal{S}} x_{g,i}$ be the occupancy of stratum $\mathcal{S}$. A stratum is *informative* for g when partially occupied,

$$\mathcal{I}_{\mathcal{S}}(x_g) = \mathbf{1}[0 < k_{\mathcal{S}}(x_g) < |\mathcal{S}|], \quad (14)$$

since only then does the within-strata shuffle have any freedom at all. The operative identifiability quantity is the variance of the locked null distribution of the gain statistic, and the smallest lower-tail p -value any observation could achieve under it:

$$\sigma_{\text{null}}^2(g) = \text{Var}_b[G(\pi_b(x_g); \tau)], \quad p_{\min}(g) = \frac{1 + \sum_{b=1}^B \mathbf{1}[G(\pi_b(x_g); \tau) \leq g_{\min}]}{1 + B}, \quad (15)$$

where $g_{\min} = \min_b G(\pi_b(x_g); \tau)$ is the smallest gain count the panel’s occupancy admits. The classification is locked:

$$\text{class}(g) = \begin{cases} \text{IMPOSSIBLE,} & \sigma_{\text{null}}(g) = 0, \\ \text{UNDERPOWERED,} & \sigma_{\text{null}}(g) > 0 \wedge p_{\min}(g) \geq 0.05, \\ \text{TESTABLE,} & p_{\min}(g) < 0.05. \end{cases} \quad (16)$$

Partial occupancy is necessary but not sufficient for identifiability: with three bdelloids, a $k = 2$ occupancy permits three arrangements, yet all can yield the same gain count because the stratum subtends one resistant branch; the null variance, not the occupancy count, is the decisive quantity.

Panel verdict: no candidate is testable. Nine candidates are IMPOSSIBLE – the structured null is a constant, so the convergence question cannot be asked of them on this panel at all –

Table 2: Per-candidate Identifiability Scores on the 20-species panel (10,000 within-strata shuffles, locked seed; all observed statistics reproduce the committed H4c values). Occupancy is *Deinococcaceae/Bdelloida*.

Candidate	occupancy	inf. strata	G_{obs}	null SD	p_{obs}	p_{min}	class
WP_012692117.1	6/6, 3/3	0	2	0	1.0	1.0	IMPOSSIBLE
UJR36058.1	6/6, 2/3	1	2	0	1.0	1.0	IMPOSSIBLE
UJR11615.1	6/6, 3/3	0	2	0	1.0	1.0	IMPOSSIBLE
UJR17759.1	3/6, 1/3	2	2	0.458	0.299	0.299	UNDERPOWERED
WP_272975990.1	6/6, 3/3	0	2	0	1.0	1.0	IMPOSSIBLE
UJR28130.1	6/6, 3/3	0	2	0	1.0	1.0	IMPOSSIBLE
UJR26281.1	6/6, 3/3	0	2	0	1.0	1.0	IMPOSSIBLE
WP_012692142.1	6/6, 3/3	0	2	0	1.0	1.0	IMPOSSIBLE
WP_414646549.1	6/6, 2/3	1	2	0	1.0	1.0	IMPOSSIBLE
WP_272975981.1	6/6, 3/3	0	2	0	1.0	1.0	IMPOSSIBLE

and the single candidate with real null variance is UNDERPOWERED: even its most extreme achievable configuration yields $p_{\text{min}} = 0.299$, an order of magnitude from significance. This sharpens the H4c statement (“power exists only for partially occupied strata”): partial occupancy is where power *could* live, but this panel’s one partially occupied candidate shows it need not live there. The score thereby converts a qualitative limitation into a reusable diagnostic with an explicit remedy: only phylogenetically *dispersed* resistant lineages raise p_{min} below the decision threshold – additional proteomes within the same clades change nothing, which is the quantitative content of the identifiability boundary. On the original 15-species panel, all ten candidates were IMPOSSIBLE (committed H4: zero-variance nulls throughout). No H4 gate, label, or caveat changes; the score is the framework’s self-diagnostic layer, and it is what a reviewer can apply to any future panel before believing a convergence claim from it.

4.6 Genome-wide permutation FDR of the candidate screen (amendment h8b, 2026-09-27 12:13 IST, locked before computation)

The empirical- p gate calibrates each orthogroup separately; it does not by itself bound the genome-wide false discovery rate of the twelve-candidate list. We estimate that rate by replaying the entire locked decision rule on every one of the screen’s own 1000 label permutations. The screen was re-run from its committed implementation with the full statistic matrix retained, $S \in \mathbb{R}^{52,786 \times 1001}$, where $S_{g,0}$ is the observed one-sided Fisher statistic of orthogroup g and $S_{g,b}$ its value under permutation π_b (same input, arms, hypergeometric kernel, seed, and permutation count as the H1 gate; a pre-committed sanity lock reproduced all twelve committed hit identities with their exact stored empirical p -values - eleven at 13/1001, one at 17/1001 - before any FDR quantity was computed). Each permuted labeling is then scored by the identical rule, with its empirical p obtained by rank recursion against the same matrix:

$$N(\pi_b) = \# \left\{ g : \underbrace{\frac{1 + \sum_{j \neq b} \mathbf{1}[S_{g,j} \leq S_{g,b}]}{1001}}_{\text{rank-recursion empirical } p} < 0.05 \wedge p_g^R(b) > p_g^S(b) \wedge \sum_{o=1}^3 \text{origin}_{g,o}(b) \geq 2 \right\}, \quad (17)$$

and $\widehat{\text{FDR}} = \frac{1}{12B} \sum_{b=1}^B N(\pi_b)$. The result is honestly negative for genome-wide control: permuted labelings yield $N(\pi_b)$ with mean 22.3, median 0, and maximum 2,561, giving $\widehat{\text{FDR}} \approx 1.85$ (normal 95% upper bound 2.66). The null is strongly bimodal - most permuted labelings produce no hits at all, but labelings that preserve phylogenetic correlation structure generate hit sets larger than the observed one. Two qualifications are recorded rather than tuned away: the estimator scores each permuted labeling against the fixed permutation set (exact only in

expectation under redrawn permutations), and the mean is inflated by correlation-preserving labelings that the ≥ 2 -origins gate only partially removes. The consequence is stated plainly: the twelve candidates carry the strongest per-orthogroup evidence this panel can produce, but no genome-wide error-rate control is claimed for the list; confirmation routes through the mechanistic prioritization arm and experiment. An earlier approximation of this analysis (raw hypergeometric Fisher on a presence-only matrix) aborted on its own pre-committed sanity lock when it failed to reproduce the committed hits; the abort and its root cause are archived in `results/h8_genomewide_fdr.json`.

4.7 Baselines: simpler selection rules against the framework (amendment 2026-09-27 12:19 IST, locked before list generation)

Three simpler selection rules were executed against the committed twelve under a locked arm-blind grading protocol (pool shuffled with a fixed seed before any grading; the H5 v2 tier definitions applied unchanged). B0 RANDOM: uniform draws from the tested universe. B1 SEQUENCE-ONLY: mmseqs2 search of the panel proteomes against the locked S^+ reference families, top twelve by e -value. B2 DOMAIN-ONLY: per-orthogroup resistant-direction Fisher on a panel-wide hmmscan of the locked 64-family Pfam set, with no specificity and no origins layers. Grading outcome: B0 yields 0/36 Tier-1-or-2 entries; B1 yields 12/12 Tier-1; B2 yields 8/12 Tier-1-or-2; the framework (B3) yields 8/12. Under the locked strict-exceed rule the framework beats random selection, does not beat sequence-only selection, and ties domain-only selection - and both non-wins are structurally informative. B1's perfect Tier-1 score is tautological: it selects by sequence similarity to literature-known families, so all twelve of its calls are the canonical families themselves, including copies in *sensitive* species (a *Thermus* DdrA homolog, *Hypsibius* CAHS/SAHS) - it rediscovers the known and finds nothing new, which is precisely what an enrichment screen must look past. B2 returns a list *identical* to the committed twelve: on the annotatable universe the specificity and origins layers confirm rather than filter, evidence that the framework's added machinery costs no candidates. The comparison thus converts a reviewer question ("would something simpler do?") into a measured answer: simpler rules either rediscover known families or recover the same list with less audit structure.

4.8 Sensitivity audits: orthogroup calling and assembly quality (amendments 2026-09-27 12:29 IST, locked before computation)

Two robustness audits were executed under locked report-only rules. *Orthogroup-calling stability*: re-mapping the committed hits' resistant-arm members across the two alternative clusterings on disk, 11 of 12 hits co-cluster ($\geq 80\%$ of resistant members) in the 13-species mmseqs v2 clustering; the single splitting hit (UJR10666.1, folate-pathway) is flagged here and in Caveats rather than removed. Co-clustering under linclust30 is low (3/12), which matches linclust's documented fragmentation on this data (subset concordance purity 0.914, retention 0.363) and is why formal claims never rest on linclust counts. *Assembly-bias control*: hit-presence counts per genome show no correlation with proteome size (Spearman $\rho = -0.095$, $p = 0.78$), resistant and sensitive arms do not differ in proteome size (medians 3,479 vs 12,974; MWU $p = 0.93$, small- n disclosed), and no hit concentrates in outlier-size genomes (0/12 fail the locked $2\times$ -median check). The committed twelve are not an assembly-quality artifact under either audit.

4.9 Radiation-specific or general stress? A classification power limit (amendment 2026-09-27 12:38 IST, locked before computation; delegated execution, reviewed and committed by the lane)

To ask whether the committed candidates reflect radiation-specific or general stress adaptation, the H1 gate was re-run under two alternative sensitive-arm compositions: lab-model controls only (S-A) and stress-specialist controls only (S-B), both at $N = 8$. All twelve candidates classify CONTROL-DEPENDENT: none passes the gate under either composition (empirical $p = 0.075\text{--}0.081$). The sweep is a power property, not biology: every candidate carries zero sensitive copies under both compositions, making S-A and S-B numerically identical, and at $N = 8$ with $p_s = 0$ the permutation floor is $\binom{5}{p_m}/\binom{8}{p_m} - 0.071$ at $p_m = 4$, unreachable by any candidate at the locked 0.05 gate, while the same $p_m = 4$ profile scores 0.0152 under the committed $N = 11$ panel. The classification question is therefore *untestable at reduced panel size* (the identifiability theme of Section 4.5 recurring at the level of arm composition); separating radiation-specific from general-stress adaptation needs more independent sensitive lineages, not a re-run. Recorded verbatim; no candidate labels change.

4.10 Functional evidence tiers (queue #5, amendment 2026-09-27 13:25 IST)

Each of the 12 candidates was classified against a four-tier functional-evidence ladder using committed data only: T1 domain presence (11/12 pass; UJR09360.1 has no committed domain record), T2 orthology confidence (11/12 pass), T3 expression (not available - no committed transcriptomes), T4 phenotype association (not available). Every candidate's evidence stops at T2. The tier table makes the evidential ceiling explicit per candidate: domain names (ankyrin repeats, DHH/DHHA2, DSBA/thioredoxin, HAD hydrolase, flavin reductase, acetyltransferases, and others, verbatim in results/h_functional_tiers.json) describe architecture, not demonstrated function, and no candidate carries expression or phenotype evidence. The candidate list is therefore published as a tier-annotated resource, not as confirmed radiation-specific machinery.

4.11 Ranked public resource: stress-adaptation genomic atlas v1 (queue #17)

The 12 candidates are published as a tier-annotated, ranked atlas (data/stress_atlas_v1.json; Table 3). The ranking key was locked before assembly. The atlas is a resource for follow-up, not a confirmed-gene list: every entry carries its tier ceiling (all stop at T2), its CONTROL-DEPENDENT specificity class (h12), and the genome-wide FDR context (h8b, $\text{FDR} > 1$).

4.12 Blind known-gene recovery (queue #6, amendment 14:12 IST)

With S^+ family labels held out, the observed screen statistic was recomputed for all 52,786 orthogroups and ranked. Zero of the 12 unique S^+ orthogroups passes the screen gate or ranks in the top 100 - an honest negative - but the four *Deinococcus*-specific repair genes (PprI, DdrB, PprA, Dps1) rank 290–429 (median 357, top $\sim 0.7\%$), a ~ 75 -fold top-rank enrichment over uniform expectation. The gate's rejection of known stress genes is mechanistic, not statistical: universal stress machinery is shared with sensitive controls (DdrA's best hit is a *Thermus* protein), and the locked gate selects resistant-*specific* retention. The validation therefore separates

Table 3: Stress-adaptation genomic atlas v1 (queue #17): the 12 committed candidates, ranked by screen strength, with domain architecture, evidence tier ceiling, and specificity class. Every cell traces to a committed result file.

Rank	Candidate	Domains (h3)	Tier ceiling
1	Adineta_vagalUJR10666.1	HPPK, Pterin_bind	T2
2	Deinococcus_desertiiWP_012692117.1	Ank, Ank_2, Ank_3	T2
3	Adineta_vagalUJR36058.1	DUF4385	T2
4	Adineta_vagalUJR11615.1	DHH, DHHA2	T2
5	Adineta_vagalUJR17759.1	Amidase, Bac_rhamnosid, Bac_rhamnosid6H	T2
6	Deinococcus_geothermalislWP_272975990.1	DSBA, Thioredoxin_4	T2
7	Adineta_vagalUJR28130.1	3Beta_HSD, Epimerase, GDP_Man_Dehyd	T2
8	Adineta_vagalUJR26281.1	HAD_2, Hydrolase, Hydrolase_like	T2
9	Deinococcus_desertiiWP_012692142.1	Flavin_Reduct	T2
10	Deinococcus_geothermalislWP_414646549.1	PAD_porph	T2
11	Deinococcus_geothermalislWP_272975981.1	Acetyltransf_1, Acetyltransf_10, Acetyltransf_7	T2
12	Adineta_vagalUJR09360.1	none committed	T0

two questions the queue item compressed: the statistic detects stress biology (enriched), and the gate trades known-gene recovery for novelty (by design, now quantified).

4.13 Independent-clade confirmation: *Deinococcus indicus* (queue #11, amendment 14:18 IST)

A sixth independent confirmation proteome (*D. indicus* GCA_014653295.1, 4,153 proteins, sha256-ledgered) was scanned under the committed h3 rule (full v1 domain set in ≥ 1 protein, $E \leq 10^{-5}$). Nine of eleven domain-annotated candidates gain *D. indicus* support; both non-supported candidates are bdelloid-originating (*Adineta vaga* UJR10666.1, UJR17759.1), exactly where a *Deinococcus* proteome should not confirm - the arm adds an independent within-family replication for the *Deinococcaceae* candidates and records the bdelloid non-confirmation verbatim. A seventh proteome (*D. wulumuqiensis* GCA_053878665.1, 3,173 proteins, amendment 14:20 IST) reproduces the support pattern exactly (9/11, identical per-candidate outcomes), so the within-family confirmation is consistent across two distinct *Deinococcaceae* clades. Both sit INSIDE origin 1 (*Deinococcaceae*): these arms replicate within-origin and do not add an independent evolutionary origin - the queue #2 arm below supplies the first true independent-origin test. Two bdelloid proteomes (*Rotaria sordida*, 44,209 proteins; *Adineta steineri*, 54,870 proteins; amendment 14:21 IST, within-origin-3 replications) complete the matrix: bdelloid-originating UJR10666.1 gains support in the congeneric *A. steineri* (its first confirmation), while UJR17759.1 fails all four new proteomes - a cross-clade negative that downgrades it to panel-specific. Seven of eleven candidates are supported in every new proteome.

True independent origins (queue #2, amendment committed 14:20:15 IST). *Kineococcus radiotolerans* (GCF_014191015.1, 4,368 proteins) and *Rubrobacter radiotolerans* (GCF_033842665.1, 3,244 proteins) - actinobacterial radioresistance, phylogenetically outside *Deinococcaceae*, Tardigrada, and Bdelloidea, hence genuinely independent origins of radioresistance - were scanned under the identical pre-registered rule (amendment locked 14:20:15 IST, before scanning; sha256-ledgered; results/h_kineococcus_radiotolerans_confirmation.json, results/h_rubrobacter_radiotolerans_confirmation.json). Each supports 7 of the 11 domain-annotated candidates. The cross-arm core - candidates sup-

ported in all six new proteomes spanning four independent radioresistant origins - is four: the bdelloid-originating pair UJR26281.1 and UJR28130.1 and the *D. geothermalis* pair WP_272975981.1 and WP_272975990.1. That two bdelloid-originating candidates confirm in actinobacterial radioresistants is a cross-phylum signal beyond the original three-clade design; UJR17759.1 remains supported in none of six (panel-specific), and UJR09360.1 is unscanned (no committed domain record). Per-candidate support across all six proteomes is committed in results/h_functional_tiers.json (T2_new_proteomes_supported_count). Tier records updated.

5 Caveats (honesty section)

An implementation audit (2026-09-27, correction amendment 00:54 IST) found that the first H1 run encoded a 12-slot panel vector over the 11-species true panel, placing *Thermus aquaticus* in the resistant arm with wrong 6-vs-6 Fisher/MWU constants. The locked gates were re-run on the corrected mapping with unchanged thresholds, seed, and permutation count (PRE-REGISTRATION.md): all 12 Fisher candidate identities and all specificity labels survive unchanged, the Fisher floor becomes $1/\binom{11}{5}$, the complete-separation empirical p becomes 0.013, and the secondary MWU arm shrinks from 53 to 26 hits (the earlier constants had inflated it). No downstream claim (H2 labels, H3 domain confirmation, H4 phylogenetic boundary) used the broken code path. The buggy outputs are archived in results/ for the audit trail. (1) Tardigrada origin is single-proteome (Milnesium terminally unavailable); the gate's " ≥ 2 origins" was met by Deinococcaceae+Bdelloidea in every passing orthogroup - no Tardigrada-supported convergence is claimed. (2) Empirical $p = 0.013$ is a 1000-permutation count estimate ($13/1001$) for complete-separation configurations at the true 5-vs-6 panel; the estimator floor is $1/1001 \approx 10^{-3}$ and the exact-enumeration Fisher floor is $1/\binom{11}{5} \approx 2.2 \times 10^{-3}$, so 0.013 means roughly a dozen permuted labelings scored as extreme as the observed - it is not a continuous measure. (3) Empirical- p calibration assumes exchangeability under label permutation; family-level phylogenetic structure is only partially randomized away. Section 3.4 quantifies this: structured-presence-permutation tests are degenerate on saturated clades, so convergence beyond relatedness is not claimed for any candidate (terminal identifiability statement). (4) 4 esummary records failed uid resolution and await retry; their Pfam calls stand via hmmscan. (5) Candidates are computational; no experimental validation is claimed. (6) UJR10666.1 splits in the v2 13-species clustering (Section 4.8); flagged, not removed, per the locked rule. (7) Genome-wide permutation FDR of the twelve-candidate list is ≈ 1.85 (Section 3.6): per-orthogroup calibration only; no multiple-testing control is claimed. (8) The v1 headline peroxiredoxin did not survive the specificity layer (Section 3.1); v1-era claims about AhpC-TSA as radiation-specific are superseded by the v2 labels.

5.1 Framing amendment from judge round 3 (2026-09-27)

Following the round-3 critique (ChatGPT thread, verdict logged in judge/round3_verdict.txt), the deliverable is reframed from a list of ten candidates to a *principled computational framework for prioritizing radiation-resistance mechanisms*: enrichment gating, stress-specificity filtering, orthogonal domain confirmation, and an explicit phylogenetic identifiability audit, in that order, each stage with a locked pass rule. The judge's identified hole - mechanistic prioritization beyond statistical control - becomes a new falsifiable arm: rank the confirmed candidates by mechanistic proximity to known radiation-survival biology (DNA repair, oxidative-stress response, protein protection) under a pre-locked scoring rubric, and test whether the top-ranked

candidate is the one with independent literature evidence of a radiation phenotype. The rubric is locked before any candidate scores are computed.

5.2 H5 rubric outcome: v1 falsified, v2 benchmark not earned (amendments 2026-09-27)

Under the strict locked reading, rubric v1 FAILED: the top-ranked candidate (an *A. vaga* NmrA-family redox dehydrogenase, mechanism breadth C2+C5) carries only mechanism-level evidence (3/3 PubMed family-level hits were false positives on inspection), while the first candidate with family-level published resistance evidence sits at rank 4 (DsbA/thioredoxin; PMIDs 23603741, 29518000). The failure was recorded, not patched. A v2 protocol was then locked by dated amendment before any re-evaluation (judge redirect, option (c)): the deliverable is the ranked shortlist with a two-tier evidence classification (Tier 1 family-level literature; Tier 2 mechanism-level; Tier 0 neither), benchmarked against a locked reference standard. After a pre-evaluation correction (the canonical reference families are excluded from the candidate pool by the screen's own sensitive-control filter, so enrichment was re-anchored to rank-blind Tier-1 grading), the v2 evaluation gives: 1 Tier-1, 6 Tier-2, 4 Tier-0 among the 11 candidates; P@5 Tier-1-or-2 = 5/5 but no Tier-1 family in the top 3; hypergeometric enrichment $p = 0.4545$ (with a single Tier-1 family in the pool the test floor is 0.4545 - an identifiability/power limit, recorded rather than tuned); MRR 0.25, EF@5 2.2. The v2 win condition is therefore FALSE on both legs. What stands: the evidence-graded shortlist itself as the toolkit deliverable, and the explicit finding that mechanistic breadth does not rank the literature-anchored family first - the rubric rewards mechanism coverage, while published resistance biology concentrates in a chaperone/oxidoreductase axis the rubric scored lower.

6 Next experiments (falsifiable)

- (1) Heterologous expression of the AhpC-TSA orthogroup in *E. coli*: gamma-survival assay.
- (2) Tardigrade proteome addition when a second assembly appears; re-run with identical locked gates.
- (3) Domain-level convergence test (hierarchy step 2 of the design amendment) as the orthogonal confirmation.

A Machine-readable reproducibility release (queue item 12)

Every number in this paper regenerates from the repository state at the reported commit hashes. The release contains: (i) **inputs** - the 21 panel proteomes (NCBI accessions ledgered with sha256 in `data/panel_accessions.json`), the 64 Pfam HMMs (`data/pfam_accessions.json`), the 77 candidate-cluster ortholog sequences and the 9-family S^+ reference set (`data/candidate_orthologs_data/splus_reference/splus_ledger.json`); (ii) **orthogroup tables** - `results/orthology/clusters_` (primary), `clusters_v2.tsv`, `clusters_linclust30.tsv` (sensitivity arms); (iii) **trees and strata** - `data/species_tree.nwk`, `data/h4_strata.json`; (iv) **seeds and constants** - H1 seed 11, NPERM 1000; H4c seed 260926, 10,000 shuffles; h8b seed 11 with the $52,786 \times 1001$ statistic matrix; B0 seed 26092707; grading-pool shuffle 26092708; h12 seed 260927; (v) **compute scripts** - `analysis/formal_tests.py` (H1 screen), `analysis/h8b_genomewide_fdr.py`, `src/h4b_extended.py`, `src/h12_stress_specificity.py`, `analysis/h9_b2_domain_only.py`; (vi) **outputs** - every `results/*.json` cited in the text, including archived buggy outputs

(*_buggy_20260926.json) for the audit trail; and (vii) **the pre-registration ledger** - PRE-REGISTRATION.md with every amendment and erratum timestamped before its compute. Sanity locks are part of the release: h8b aborts unless all twelve committed hits reproduce exactly, and h9/h12 inputs are checked against the committed cluster table before grading or classification.